




TTFT  $\leq 1,000$  ms and request-mean ITL  $\leq 50$  ms. Cohort goodput groups tokens by request arrival hour. Cache reuse is input-token-weighted at first admission. GPU curves show exact allocation during the first 24 h; score includes full-run average GPUs and drain. These are jointly optimized configurations; curves show replicate 1, not the best fresh replicate. See JSON for all replicate scores. Prefix reuse uses RAGPulse component-ID-derived hashes. Serving metrics are simulated.